



# More Clustering Algorithms

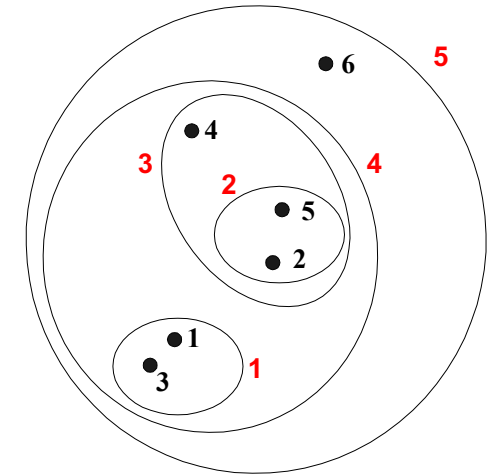
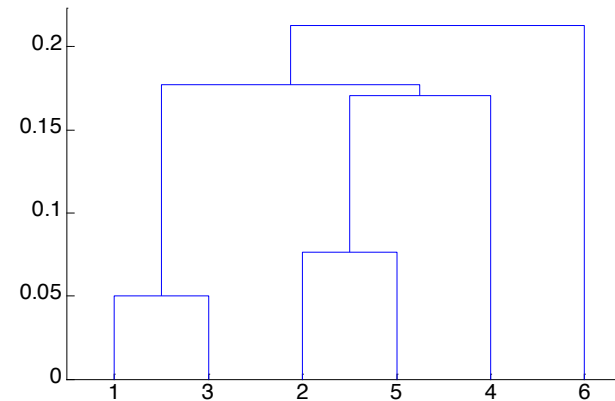
Blink Sakulkueakulsuk

# Hierarchical Clustering

Produces a set of nested clusters organized as a hierarchical tree

Can be visualized as a dendrogram

- A tree like diagram that records the sequences of merges or splits



# Strengths of Hierarchical Clustering

Do not have to assume any particular numbers of clusters

- Any desired number of clusters can be obtained by ‘cutting’ the dendrogram at the proper level

They may correspond to meaningful taxonomies

- Example in biological sciences (e.g., animal kingdom, phylogeny reconstruction, ...)

# Hierarchical Clustering

Two main types of hierarchical clustering

- Agglomerative:
  - Start with the points as individual clusters
  - At each step, merge the closest pair of clusters until only one cluster (or  $k$  clusters) left
- Divisive:
  - Start with one, all-inclusive cluster
  - At each step, split a cluster until each cluster contains an individual point (or there are  $k$  clusters)

Traditional hierarchical algorithms use a similarity or distance matrix

- Merge or split one cluster at a time

# Agglomerative Clustering Algorithm

Key Idea: Successively merge closest clusters

Basic algorithm

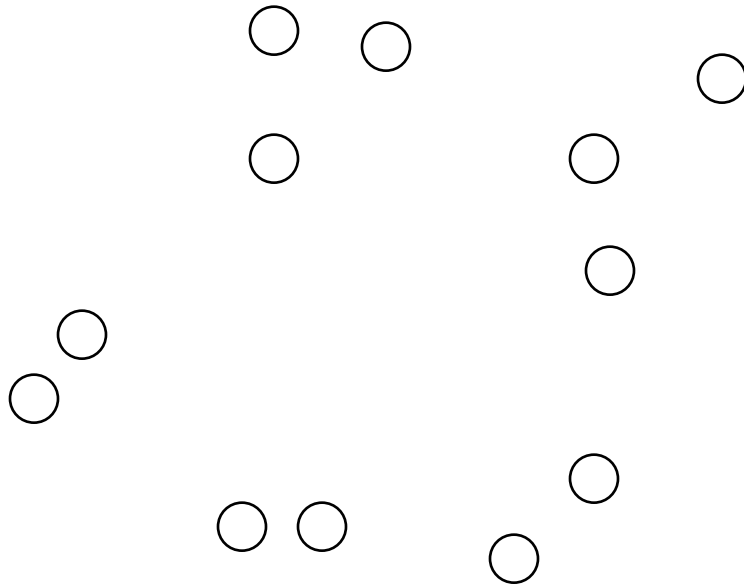
1. Compute the proximity matrix
2. Let each data point be a cluster
3. **Repeat**
4.       Merge the two closest clusters
5.       Update the proximity matrix
6. **Until** only a single cluster remains

Key operation is the computation of the proximity of two clusters

- Different approaches to defining the distance between clusters distinguish the different algorithms

# Steps 1 and 2

Start with clusters of individual points and a proximity matrix



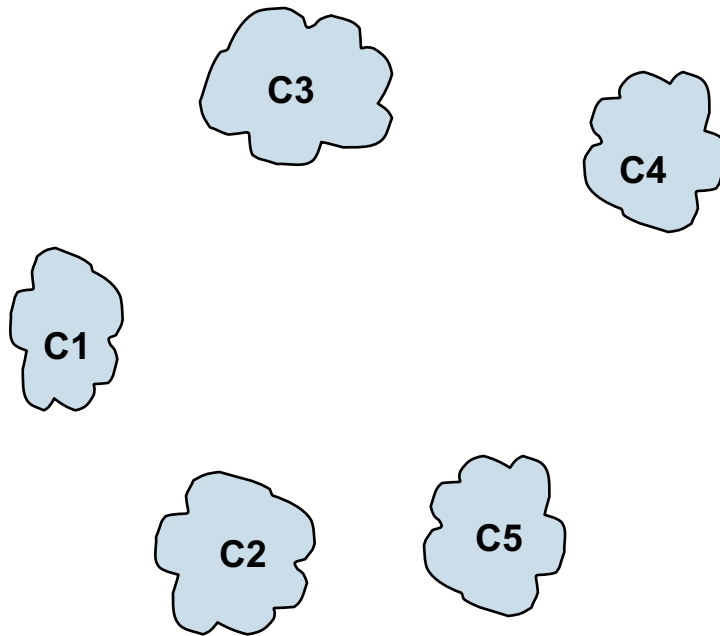
	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						

**Proximity Matrix**



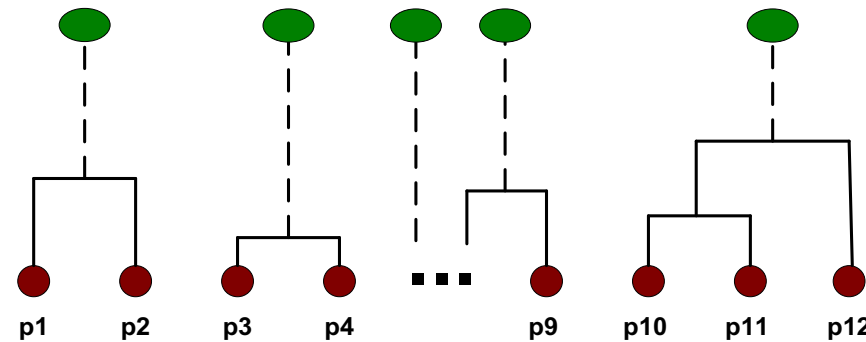
# Intermediate Situation

After some merging steps, we have some clusters



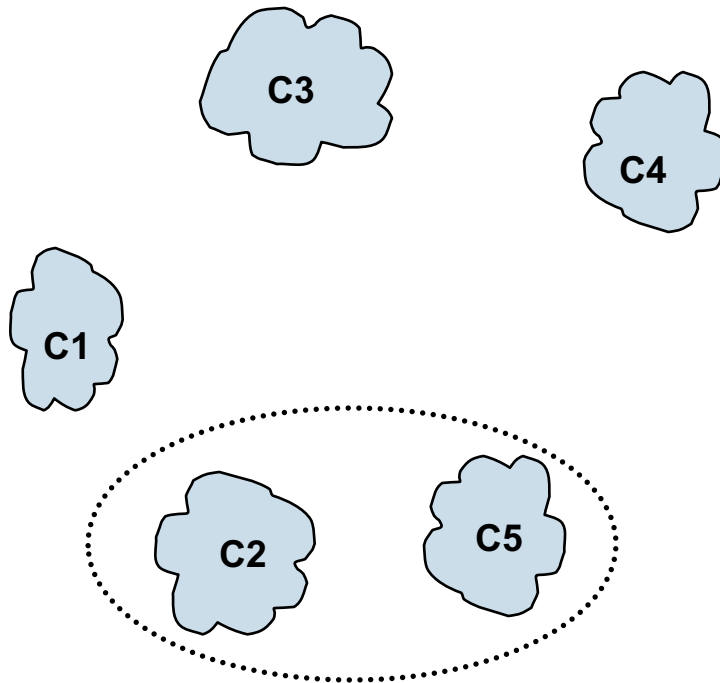
	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						

Proximity Matrix



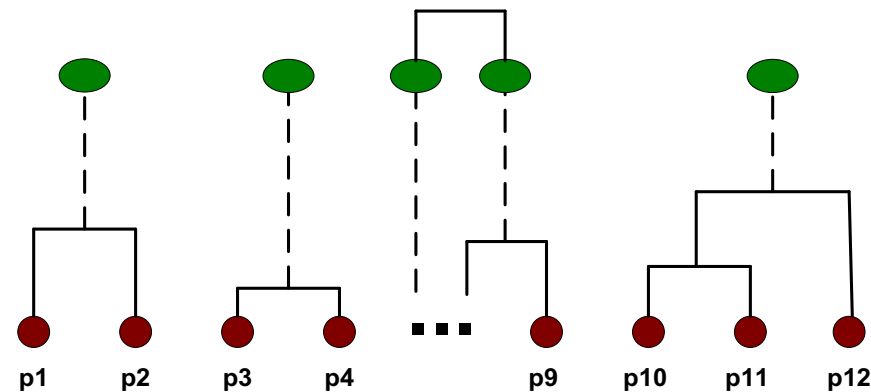
# Step 4

We want to merge the two closest clusters (C2 and C5) and update the proximity matrix.



	C1	C2	C3	C4	C5
C1					
C2					
C3					
C4					
C5					

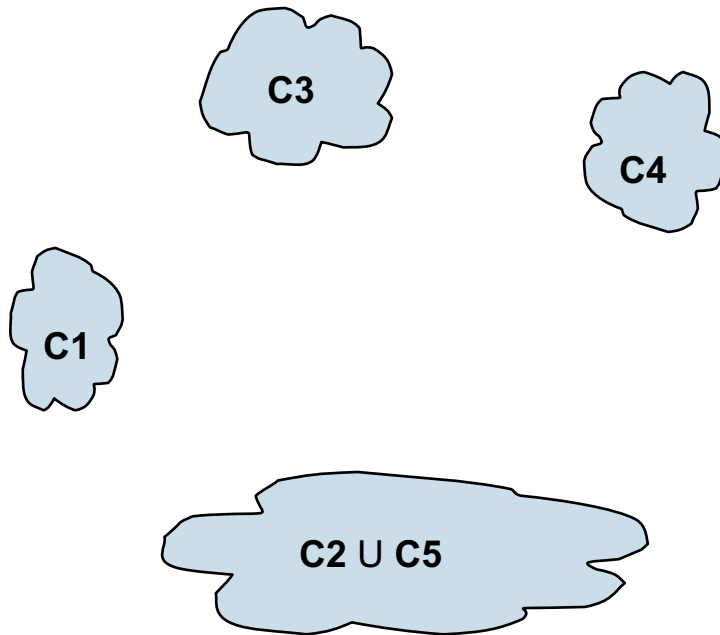
Proximity Matrix





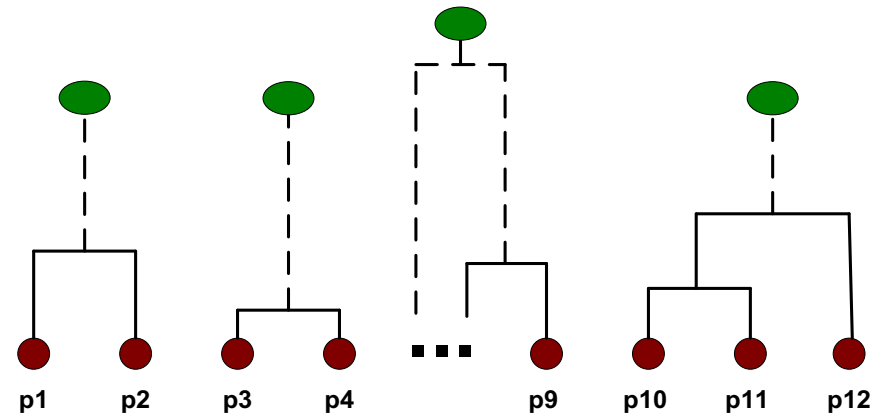
# Step 5

The question is “How do we update the proximity matrix?”

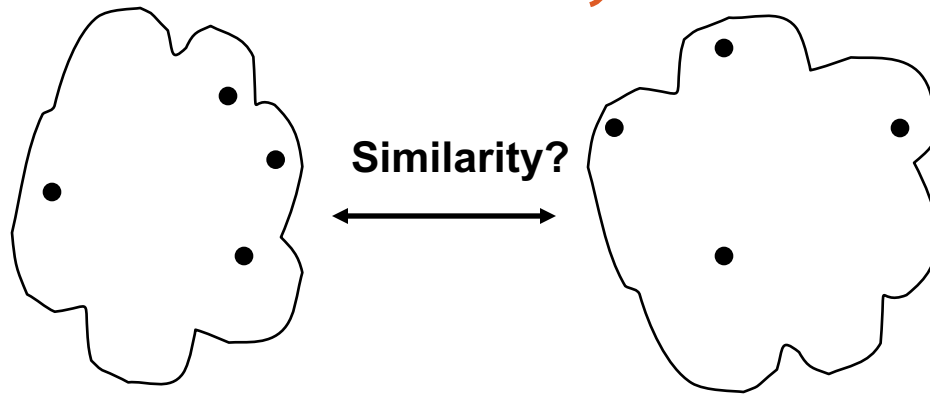


		<b>C2 U C5</b>		
	<b>C1</b>		<b>C3</b>	<b>C4</b>
<b>C1</b>		?		
<b>C2 U C5</b>	?	?	?	?
<b>C3</b>		?		
<b>C4</b>		?		

Proximity Matrix



# Defining Inter-Cluster Similarity

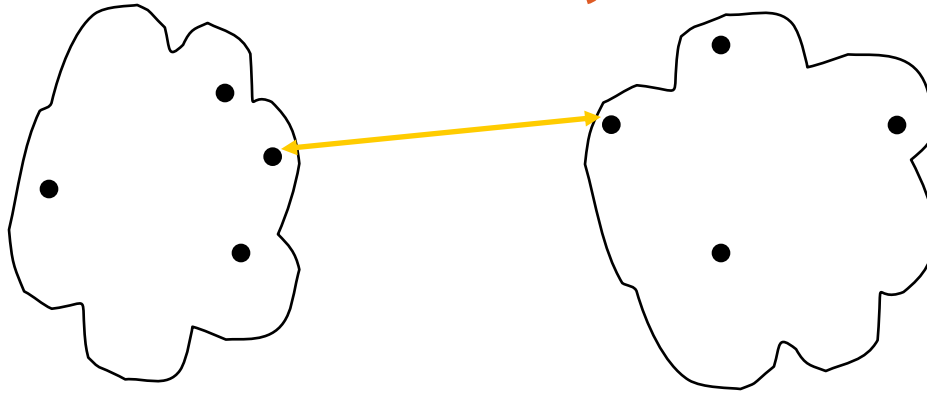


- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
  - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

**Proximity Matrix**

# Defining Inter-Cluster Similarity

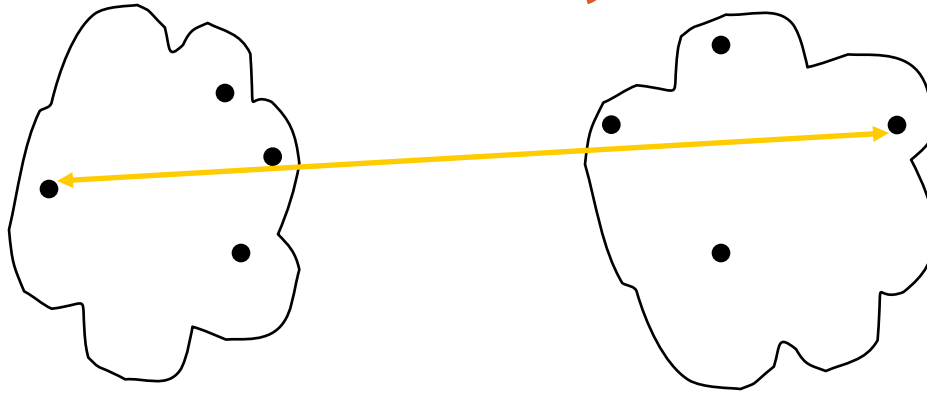


- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
  - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						

· Proximity Matrix

# Defining Inter-Cluster Similarity

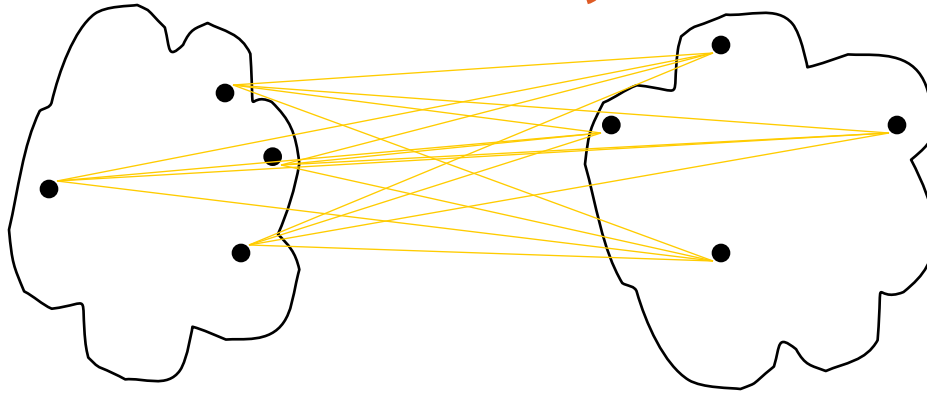


- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
  - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

**Proximity Matrix**

# Defining Inter-Cluster Similarity

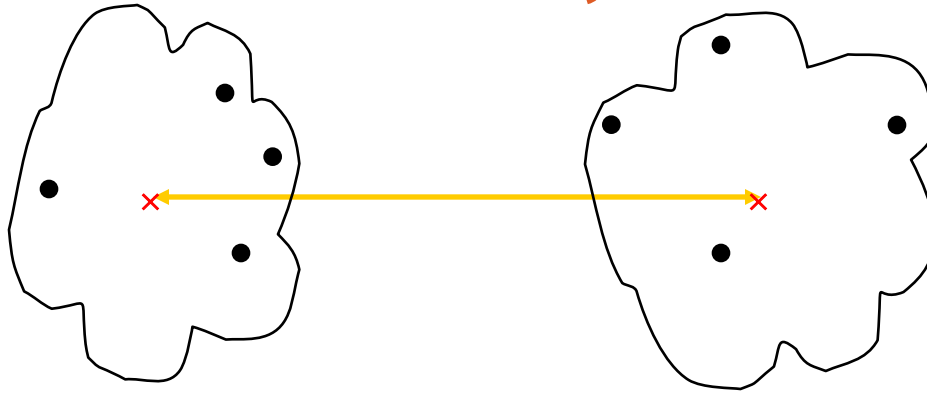


- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
  - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						

Proximity Matrix

# Defining Inter-Cluster Similarity



- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
  - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

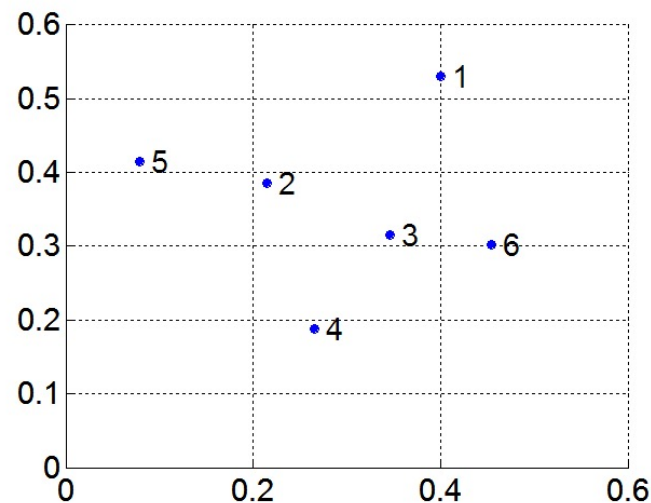
**Proximity Matrix**

# MIN or Single Link

Proximity of two clusters is based on the two closest points in the different clusters

- Determined by one pair of points, i.e., by one link in the proximity graph

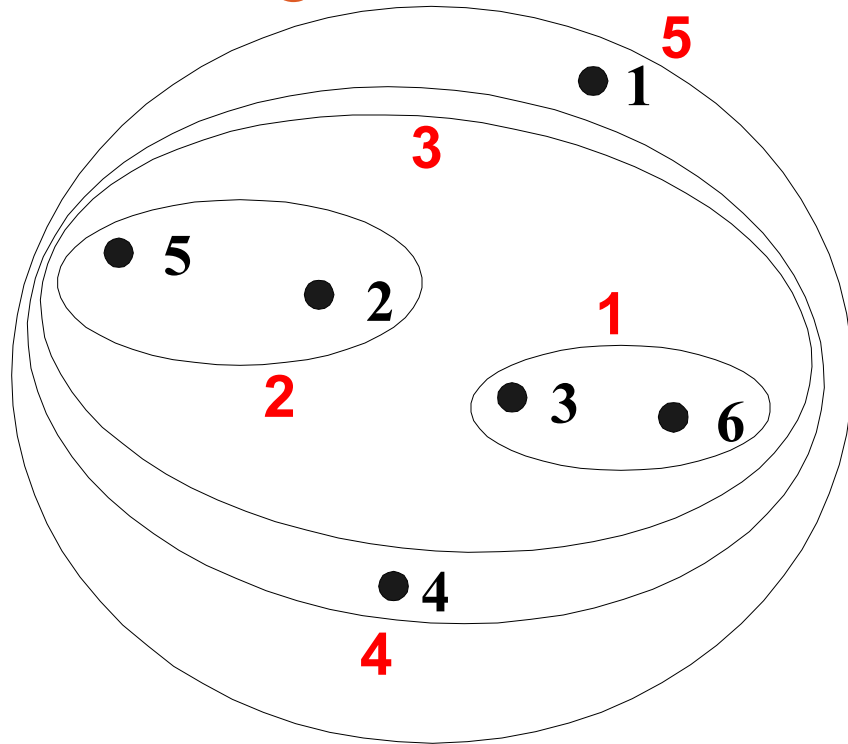
Example:



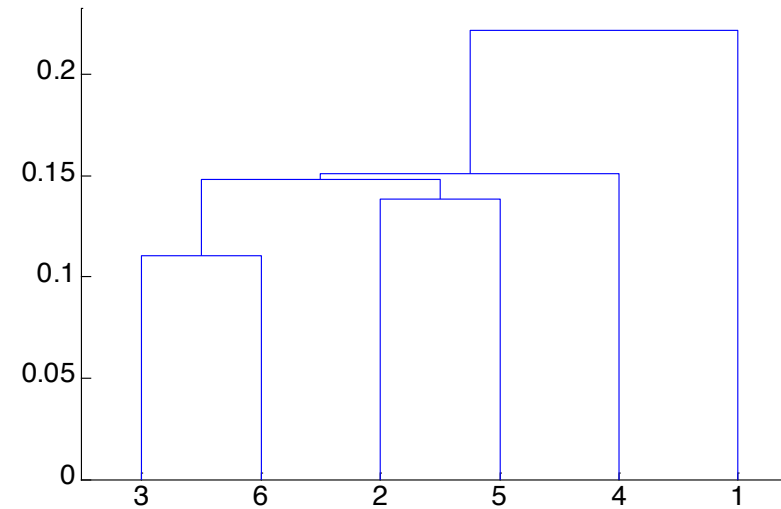
Distance Matrix:

	p1	p2	p3	p4	p5	p6
p1	0.00	0.24	0.22	0.37	0.34	0.23
p2	0.24	0.00	0.15	0.20	0.14	0.25
p3	0.22	0.15	0.00	0.15	0.28	0.11
p4	0.37	0.20	0.15	0.00	0.29	0.22
p5	0.34	0.14	0.28	0.29	0.00	0.39
p6	0.23	0.25	0.11	0.22	0.39	0.00

# MIN or Single Link



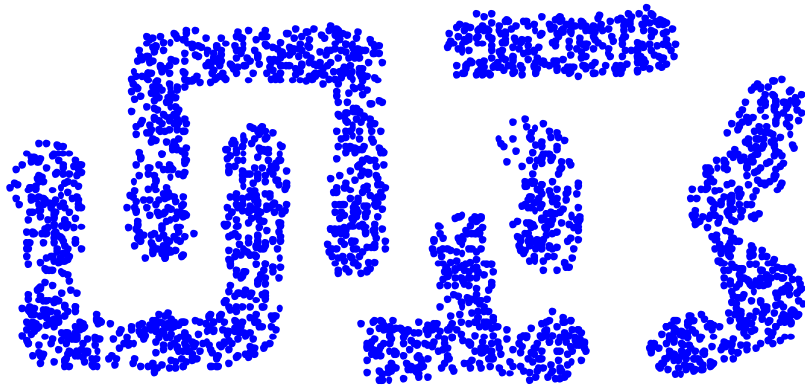
**Nested Clusters**



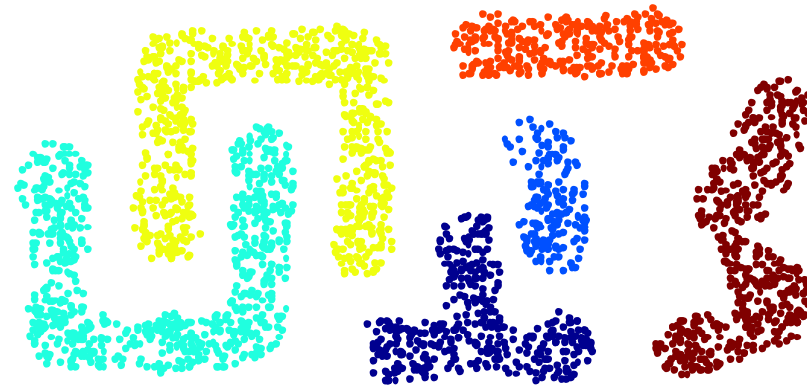
**Dendrogram**



# Strength of MIN



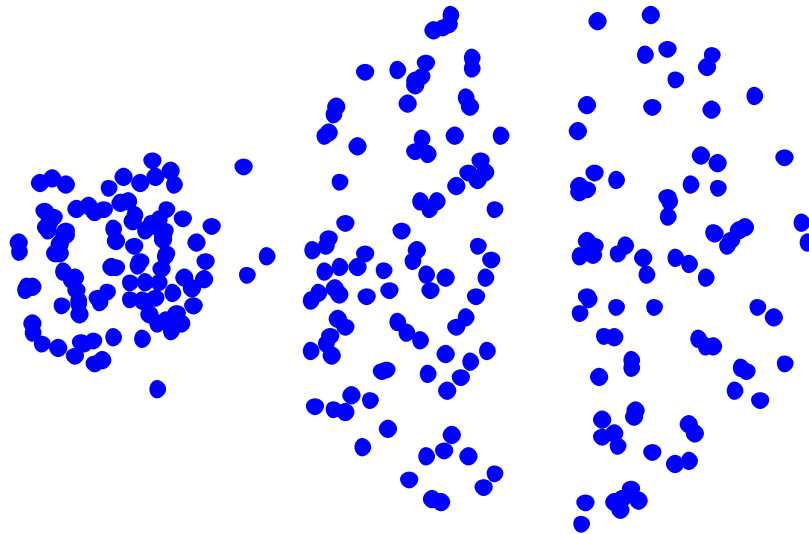
Original Points



Six Clusters

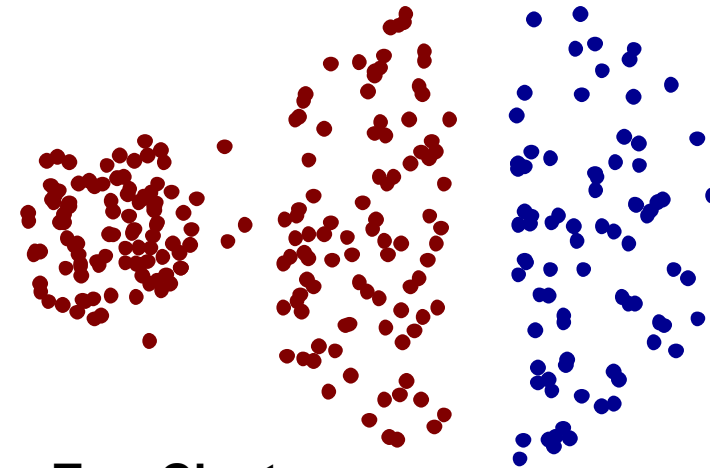
- Can handle non-elliptical shapes

# Limitations of MIN

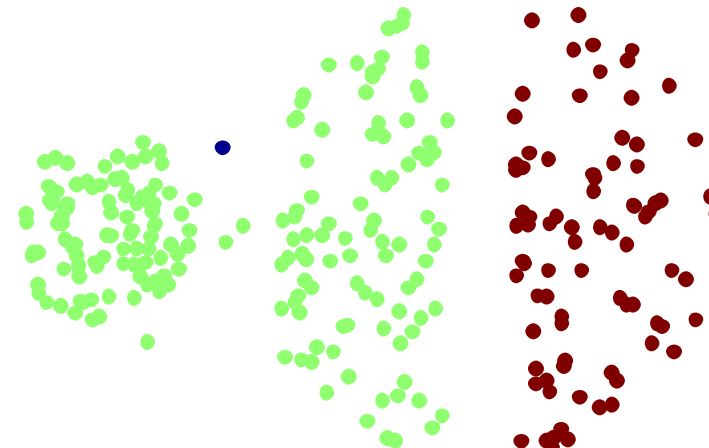


Original Points

- Sensitive to noise



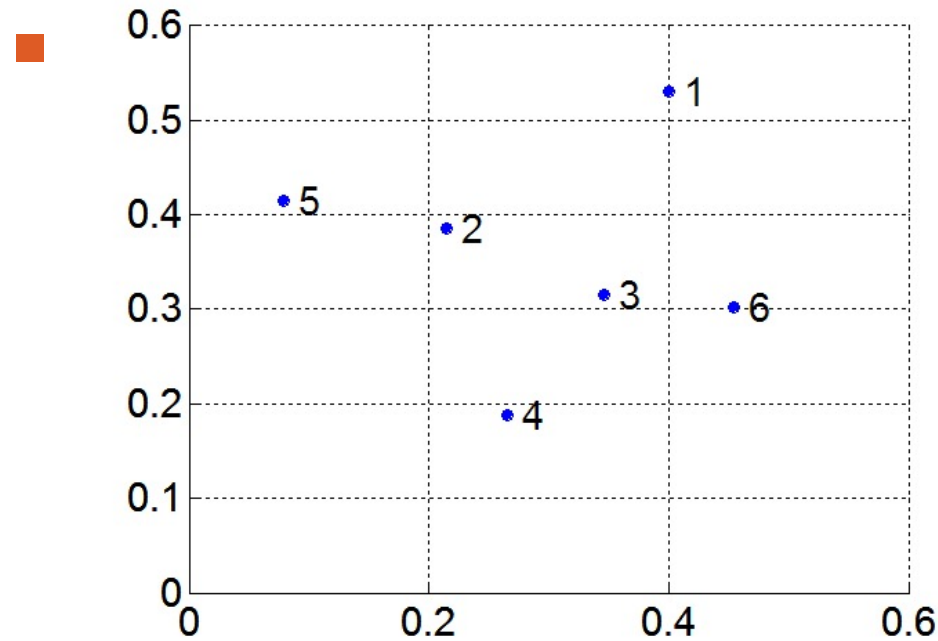
Two Clusters



Three Clusters

# MAX or Complete Linkage

Proximity of two clusters is based on the two most distant points in the different clusters

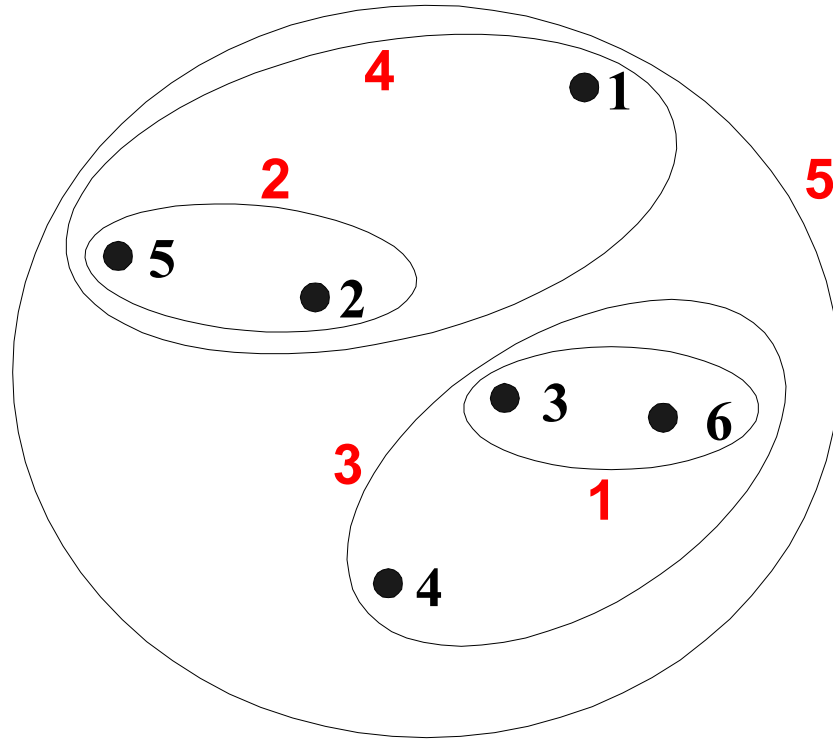


points in the two clusters

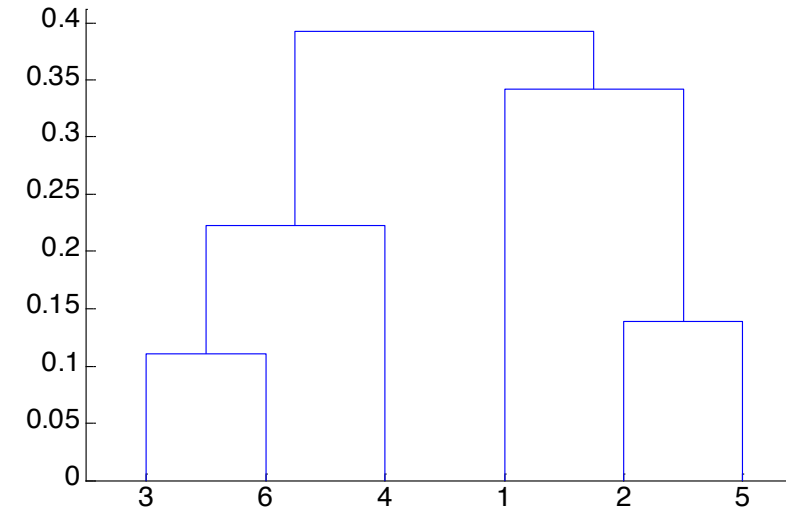
**Distance Matrix:**

	p1	p2	p3	p4	p5	p6
p1	0.00	0.24	0.22	0.37	0.34	0.23
p2	0.24	0.00	0.15	0.20	0.14	0.25
p3	0.22	0.15	0.00	0.15	0.28	0.11
p4	0.37	0.20	0.15	0.00	0.29	0.22
p5	0.34	0.14	0.28	0.29	0.00	0.39
p6	0.23	0.25	0.11	0.22	0.39	0.00

# Hierarchical Clustering: MAX

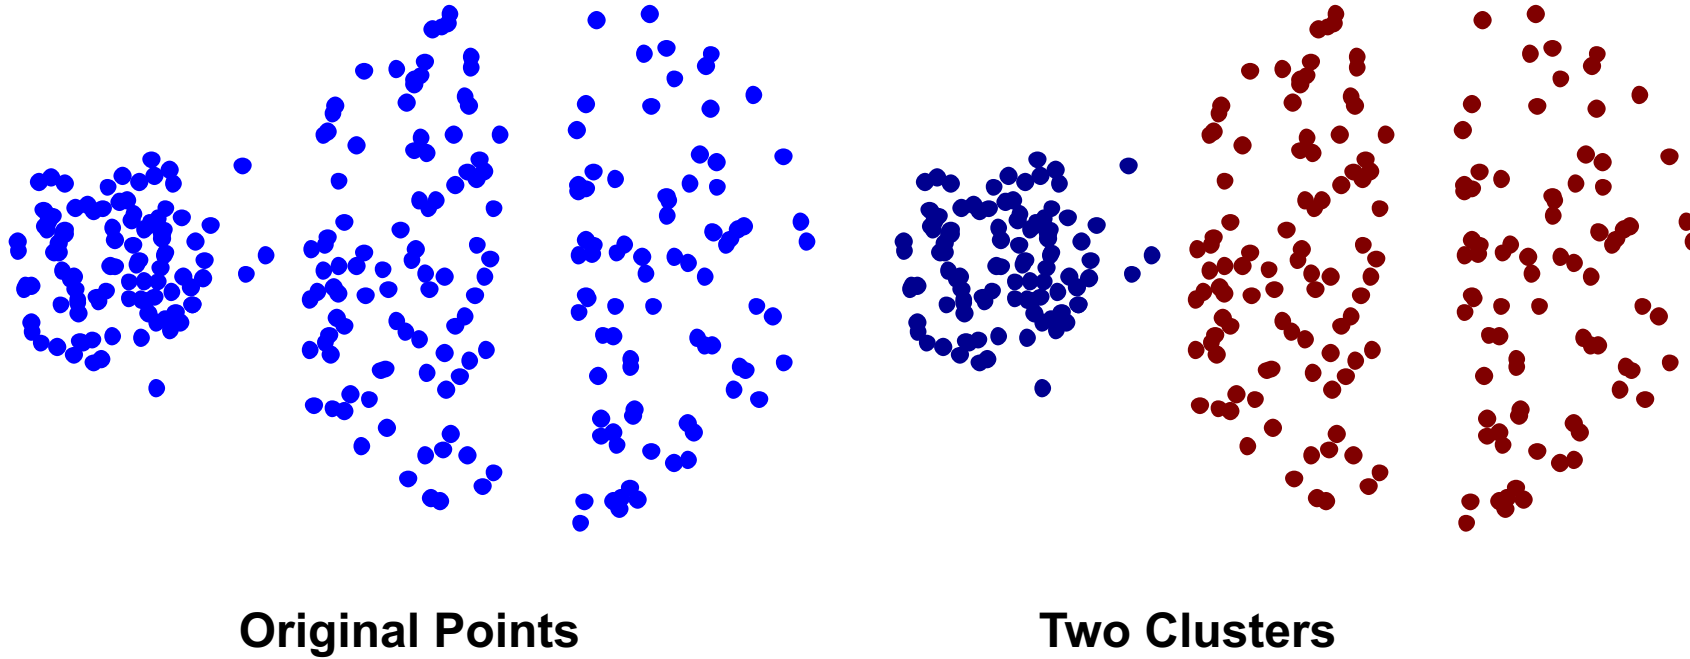


**Nested Clusters**



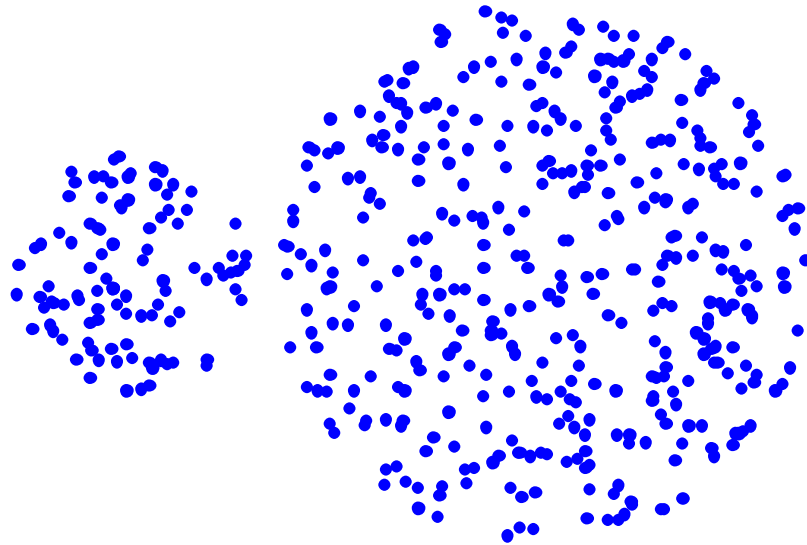
**Dendrogram**

# Strength of MAX

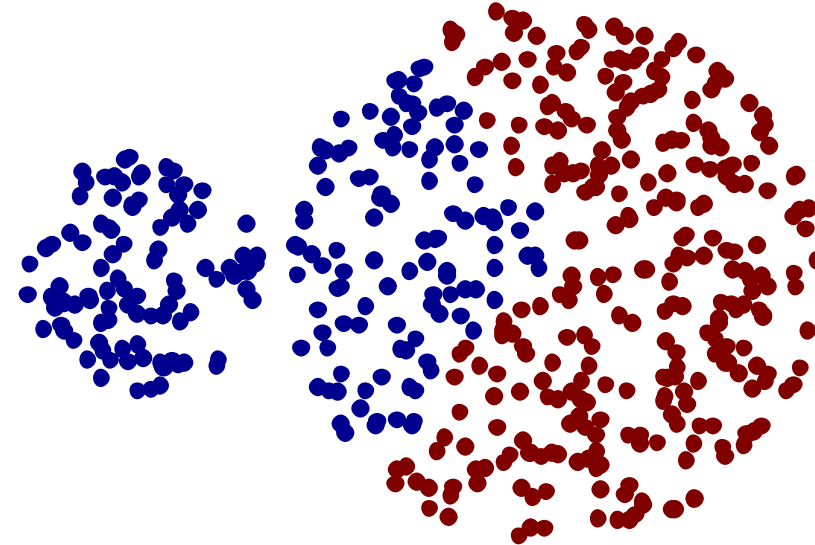


- **Less susceptible to noise**

# Limitations of MAX



Original Points



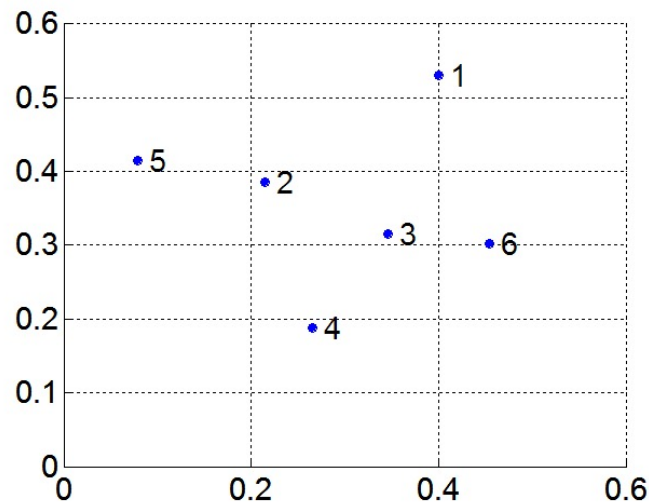
Two Clusters

- Tends to break large clusters
- Biased towards globular clusters

# Group Average

Proximity of two clusters is the average of pairwise proximity between points in the two clusters.

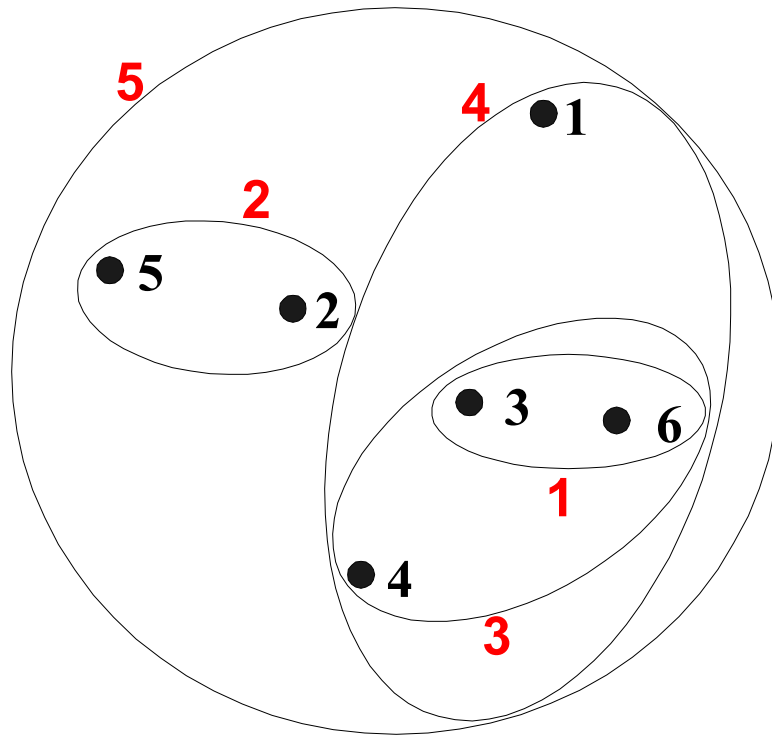
$$\text{proximity}(\text{Cluster}_i, \text{Cluster}_j) = \frac{\sum_{\substack{p_i \in \text{Cluster}_i \\ p_j \in \text{Cluster}_j}} \text{proximity}(p_i, p_j)}{|\text{Cluster}_i| \times |\text{Cluster}_j|}$$



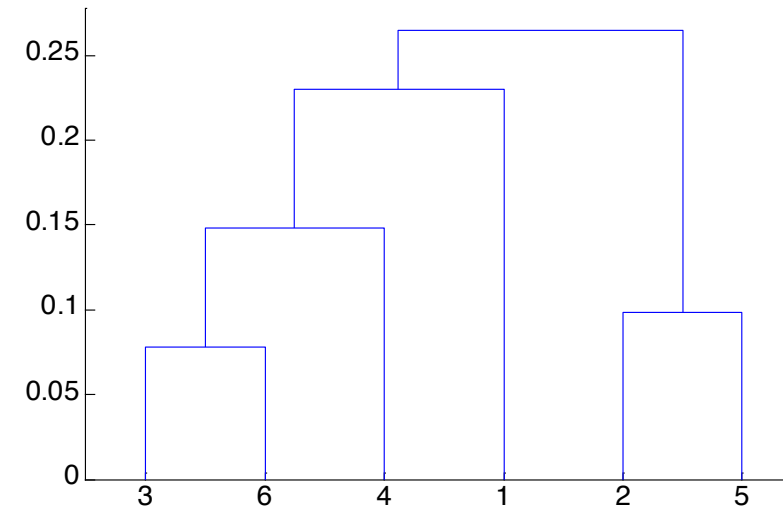
Distance Matrix:

	p1	p2	p3	p4	p5	p6
p1	0.00	0.24	0.22	0.37	0.34	0.23
p2	0.24	0.00	0.15	0.20	0.14	0.25
p3	0.22	0.15	0.00	0.15	0.28	0.11
p4	0.37	0.20	0.15	0.00	0.29	0.22
p5	0.34	0.14	0.28	0.29	0.00	0.39
p6	0.23	0.25	0.11	0.22	0.39	0.00

# Hierarchical Clustering: Group Average



**Nested Clusters**



**Dendrogram**



# Hierarchical Clustering: Group Average

Compromise between Single and Complete Link

## Strengths

- Less susceptible to noise

## Limitations

- Biased towards globular clusters

# Cluster Similarity: Ward's Method

Similarity of two clusters is based on the increase in squared error when two clusters are merged

- Similar to group average if distance between points is distance squared

$$\Delta(A, B) = \sum_{i \in A \cup B} \|\vec{x} - \vec{m}_{A \cup B}\|^2 - \sum_{i \in A} \|\vec{x} - \vec{m}_A\|^2 - \sum_{i \in B} \|\vec{x} - \vec{m}_B\|^2$$

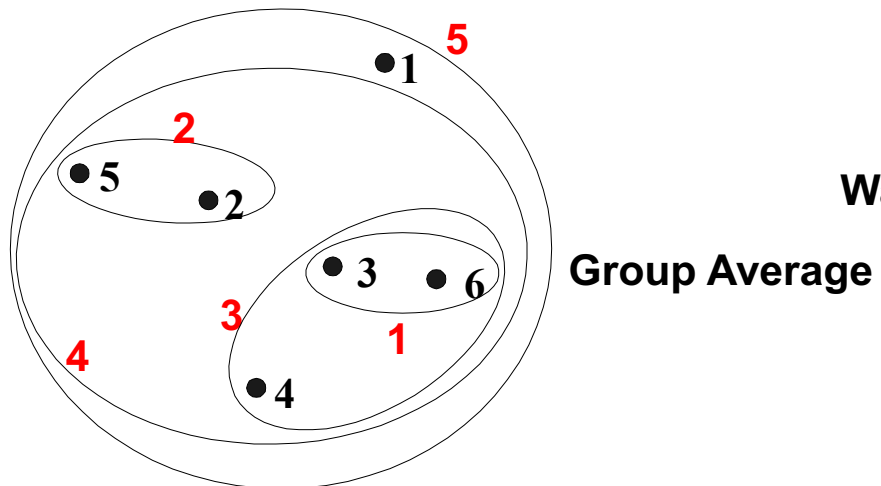
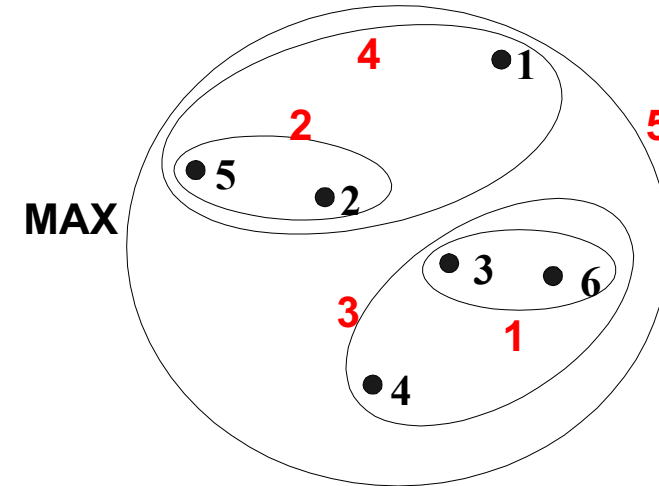
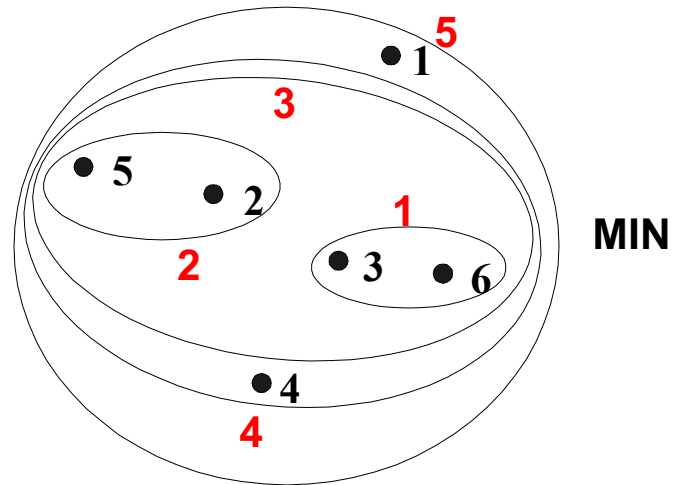
Less susceptible to noise

Biased towards globular clusters

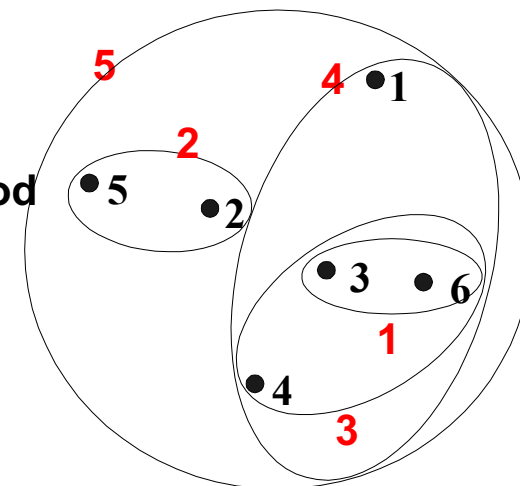
Hierarchical analogue of K-means

- Can be used to initialize K-means

# Hierarchical Clustering: Comparison



**Ward's Method**



## Hierarchical Clustering: Time and Space requirements

$O(N^2)$  space since it uses the proximity matrix.

- $N$  is the number of points.

$O(N^3)$  time in many cases

- There are  $N$  steps and at each step the size,  $N^2$ , proximity matrix must be updated and searched
- Complexity can be reduced to  $O(N^2 \log(N))$  time with some cleverness

## Hierarchical Clustering: Problems and Limitations

Once a decision is made to combine two clusters, it cannot be undone

No global objective function is directly minimized

Different schemes have problems with one or more of the following:

- Sensitivity to noise
- Difficulty handling clusters of different sizes and non-globular shapes
- Breaking large clusters

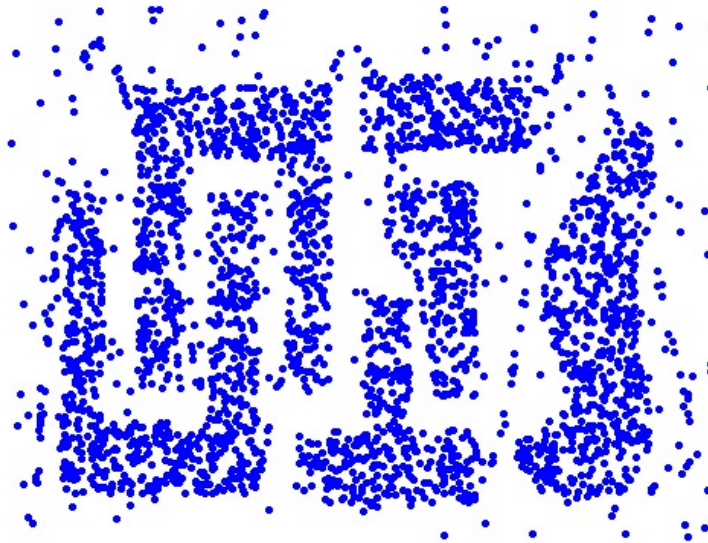
# Code Example

Performing Agglomerative Clustering and Drawing Dendrogram.

[https://scikit-learn.org/stable/auto\\_examples/cluster/plot\\_agglomerative\\_dendrogram.html#sphx-glr-auto-examples-cluster-plot-agglomerative-dendrogram-py](https://scikit-learn.org/stable/auto_examples/cluster/plot_agglomerative_dendrogram.html#sphx-glr-auto-examples-cluster-plot-agglomerative-dendrogram-py)

# Density-Based Clustering

Clusters are regions of high density that are separated from one another by regions of low density.



# DBSCAN

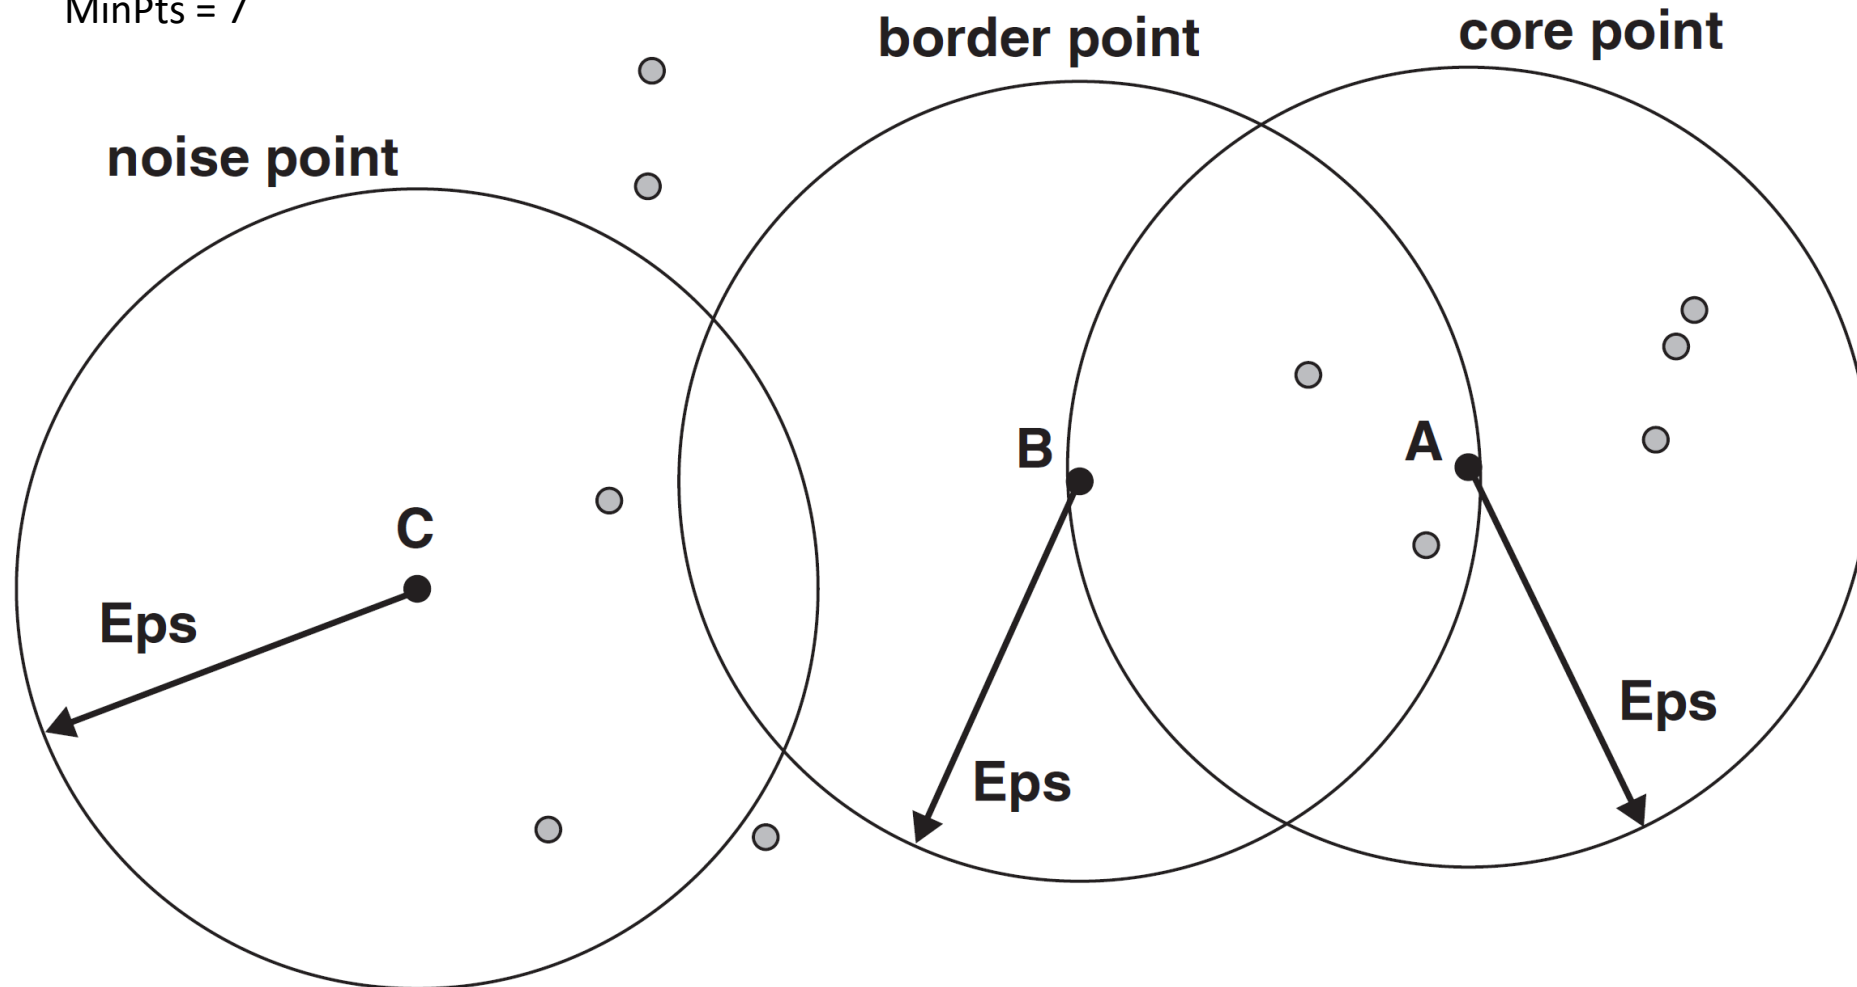
DBSCAN is a density-based algorithm.

- Density = number of points within a specified radius (Eps)
- A point is a **core point** if it has at least a specified number of points (MinPts) within Eps
  - These are points that are at the interior of a cluster
  - Counts the point itself
- A **border point** is not a core point, but is in the neighborhood of a core point
- A **noise point** is any point that is not a core point or a border point

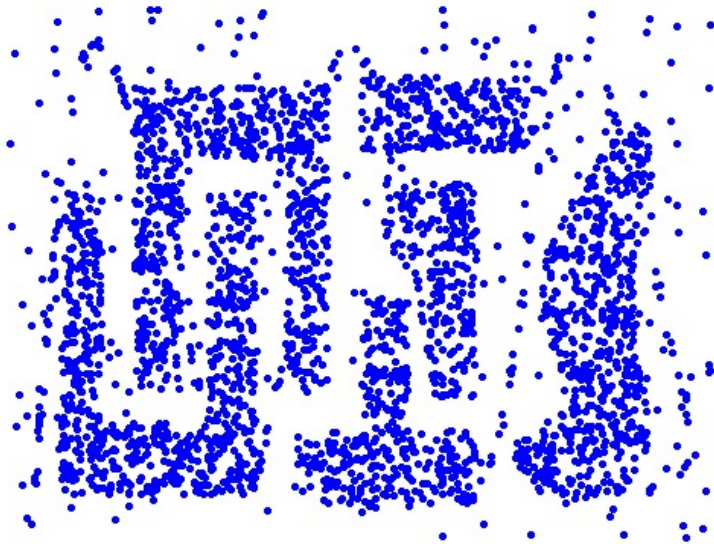


# DBSCAN: Core, Border, and Noise Points

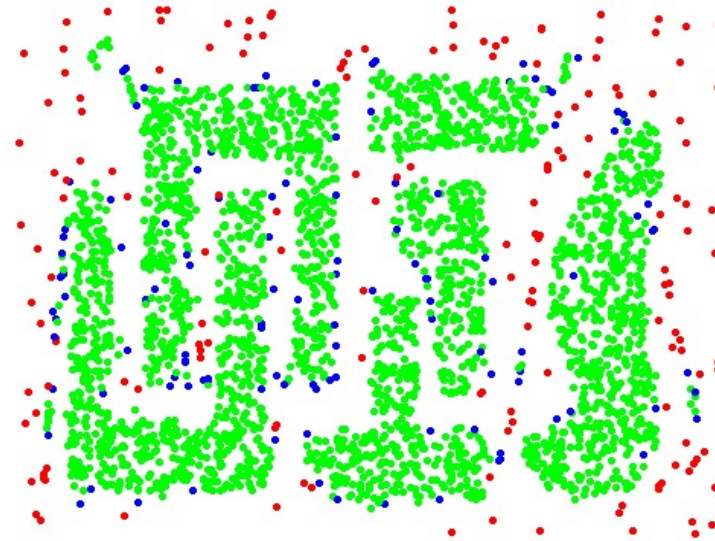
MinPts = 7



# DBSCAN: Core, Border, and Noise Points



Original Points



Point types: **core**,  
**border** and **noise**

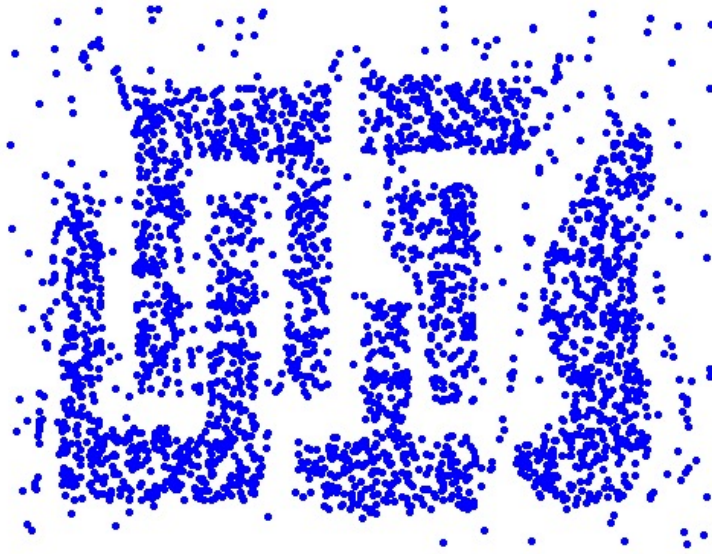
**Eps = 10, MinPts = 4**

# DBSCAN Algorithm

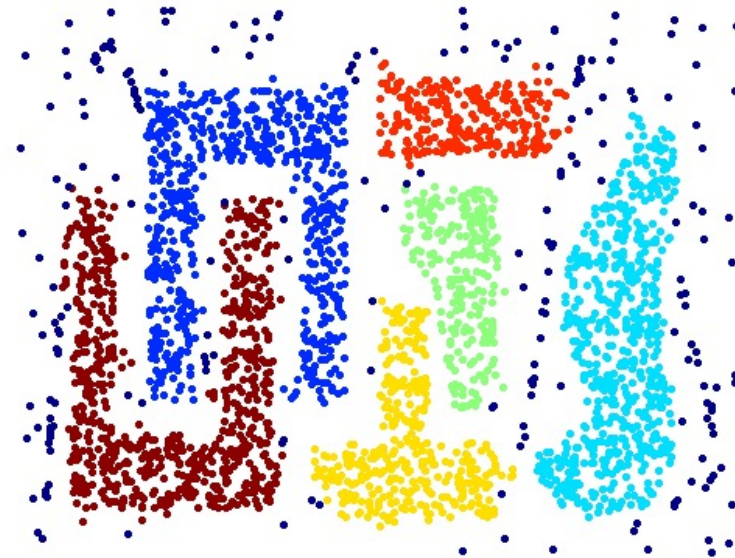
Form clusters using core points, and assign border points to one of its neighboring clusters

- 1: Label all points as core, border, or noise points.
- 2: Eliminate noise points.
- 3: Put an edge between all core points within a distance  $Eps$  of each other.
- 4: Make each group of connected core points into a separate cluster.
- 5: Assign each border point to one of the clusters of its associated core points

# When DBSCAN Works Well



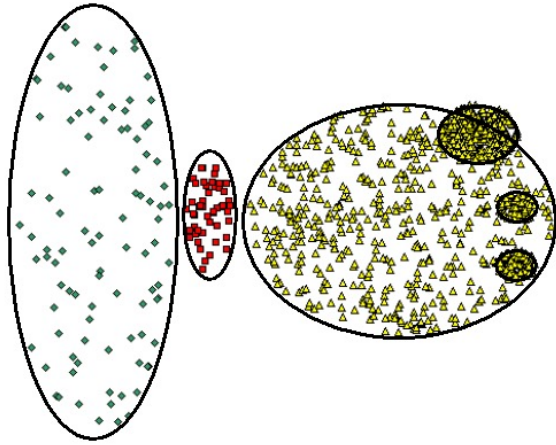
**Original Points**



**Clusters (dark blue points indicate noise)**

- **Can handle clusters of different shapes and sizes**
- **Resistant to noise**

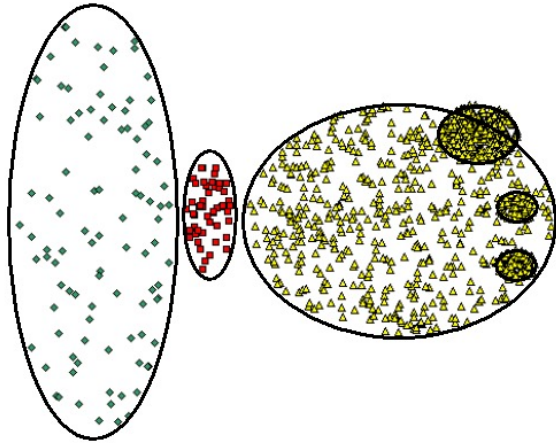
## When DBSCAN Does NOT Work Well



**Original Points**

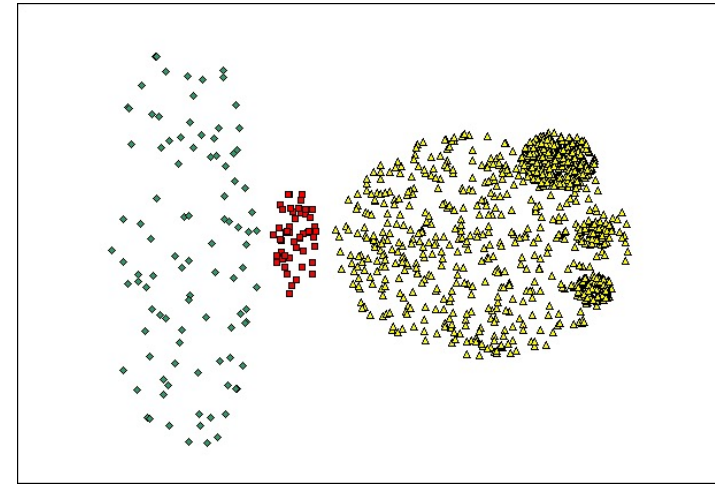


# When DBSCAN Does NOT Work Well

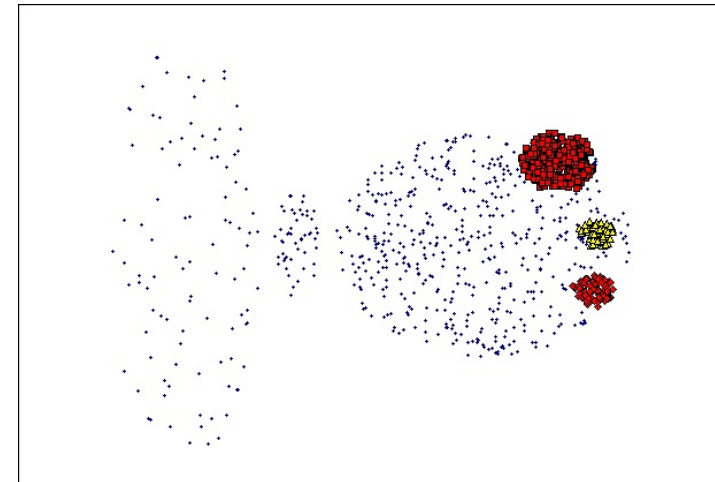


**Original Points**

- **Varying densities**
- **High-dimensional data**



(MinPts=4, Eps=9.92).



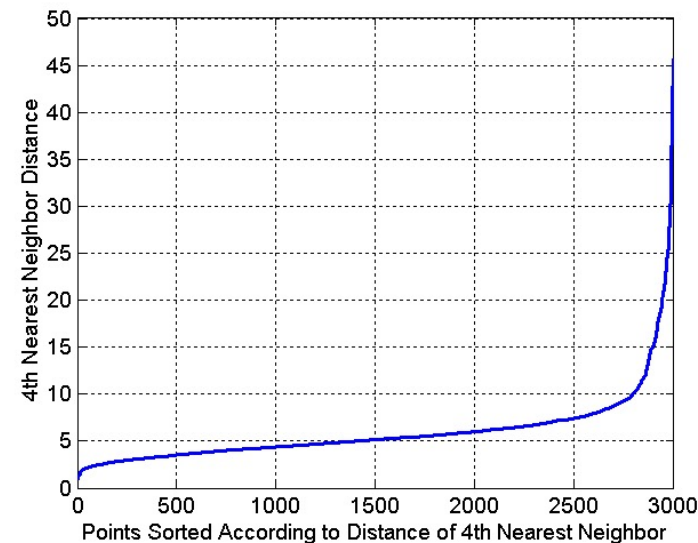
(MinPts=4, Eps=9.75)

## DBSCAN: Determining EPS and MinPts

Idea is that for points in a cluster, their  $k^{\text{th}}$  nearest neighbors are at close distance

Noise points have the  $k^{\text{th}}$  nearest neighbor at farther distance

So, plot sorted distance of every point to its  $k^{\text{th}}$  nearest neighbor



# Code Example

DBSCAN and Coloring the Clusters

[https://scikit-learn.org/stable/auto\\_examples/cluster/plot\\_dbscan.html#sphx-glr-auto-examples-cluster-plot-dbscan-py](https://scikit-learn.org/stable/auto_examples/cluster/plot_dbscan.html#sphx-glr-auto-examples-cluster-plot-dbscan-py)